

Abstract

This is the first study to compare the Regularized Horseshoe Prior (RHSP) to the Small Variance Normal Prior (SVNP) in their performance in regularizing cross loadings in Bayesian CFA. The SVNP can be used to shrink cross loadings in CFA to zero to identify models. This often results in biased model estimates, as not only small, irrelevant, but also large cross loadings are shrunk substantially. The RHSP was expected to regularize cross loadings more efficiently, avoiding the bias of the SVNP, by allowing large cross loadings to escape shrinkage within a single estimation step. It was found that indeed the SVNP had overall higher levels of bias than the RHSP with large cross loadings. Hereby, the RHSP was robust across sample sizes, and different hyperparameter settings, although under some convergence failed. Regarding the Power and Type-I-Error rate in selecting cross loadings as non-zero, both priors performed poorly, which is partially explained by the low sample sizes considered.

Keywords

Regularization, Bayesian Regularized SEM, Small Variance Normal Prior, Regularized Horseshoe Prior.

1 Introduction

The art of statistical modeling revolves around coming up with an appropriate simplification, a *model*, of a true *data-generating process*. Hereby, a fundamental trade-off between model simplicity and model complexity arises, that is mostly known as *bias-variance trade-off*. Simple models with few parameters have high bias, meaning that they deviate substantially from the true data-generating process, and low variance, such that they generalize well to other datasets from the same population. Complex models with large numbers of parameters tend to have low bias and high variance. They are thus prone to over-fitting, i.e., picking up patterns that are only relevant in the dataset at hand, but do not generalize well to other datasets. Moreover, complex models can be cumbersome to interpret and often a large number of observations is required to estimate them (Cox, 2006; James, Witten, Hastie, & Tibshirani, 2021).

In confirmatory factor analysis (CFA, see e.g. Bollen, 1989) it is common practice to deal with the bias-variance trade-off in a brute-force manner, by imposing a so-called simple structure. Here, cross loadings, factor loadings that relate items to factors that they theoretically do not belong to, are fixed to zero to yield an identified and easy to interpret model. This often leads to poor model fit, which forces researchers to free some cross loadings after the fact based on empirical grounds (modification indices) to improve fit. This procedure is flawed, as it risks capitalization on chance and thereby over-fitting (MacCallum, Roznowski, & Necowitz, 1992).

As a Bayesian solution to this issue Muthén and Asparouhov (2012) proposed identifying CFA models by setting the so-called *Small Variance Normal Prior* (SVNP) for cross loadings, which is a normal distribution with mean zero and a very small variance (e.g., $\sigma^2 = 0.01$). This prior attaches large prior mass to cross loadings of or near zero, while attaching almost no prior mass to cross loadings further from zero, such that all cross loadings in the model are shrunk. However, shrinking also the cross loadings that are

further from zero introduces substantial bias to the model (Lu, Chow, & Loken, 2016). Consequently, Bayesian CFA requires a two-step approach. First, the model is estimated with the SVNPs set for the cross loadings and cross loadings are selected as non-zero when their 95% credible intervals do not contain zero (Muthén & Asparouhov, 2012). The model is then re-estimated, where cross loadings that have been selected to be non-zero are freely estimated without shrinkage, and the remaining cross loadings are fixed to zero, avoiding the bias in the model of the previous step. Correctly selecting cross loadings as non-zero can pose a challenge in practice, as the performance of different selection criteria depends on a broad set of conditions, making it difficult to formulate general recommendations for researchers (Zhang, Pan, & Ip, 2021). Additionally, in this two-step approach, uncertainty around the selection of non-zero parameters is ignored in the final inferences based on the refitted model. This calls for shrinkage priors that allow us to ‘get a step ahead’, by regularizing cross loadings in CFA models without causing substantial bias, within a single estimation step.

One promising shrinkage prior that can be expected to allow estimating CFA models with less bias within a single step is the so-called Regularized Horseshoe Prior (RHSP), an extension of the Horseshoe Prior. This prior is designed to let large parameters escape shrinkage to zero. The RHSP has been shown to perform excellently in the selection of relevant predictors in regression (Piironen & Vehtari, 2017b; Van Erp, Oberski, & Mulder, 2019). However, systematic research on the performance of this shrinkage prior in Bayesian regularized SEM is lacking (see Van Erp, 2023 for a state of the art overview on the shrinkage priors previously used in Bayesian Regularized SEM). A notable exception is the work by Zhang and Liang (2023) who investigated the horseshoe prior in a multiple indicator multiple causes (MIMIC) model (see also Yamaguchi & Zhang, 2023 for work on the Horseshoe Prior and its extension in variable selection in a two-parameter logistic item response theory model). Empirical examples in which the performance of the RHSP in regularizing parameters in CFA models is illustrated exist (Jacobucci & Grimm, 2018a;

Van Erp, 2023). However, no previous research has validated its performance in regularizing cross loadings in CFA under a varying set of conditions with a Monte Carlo simulation study. We therefore aim to compare the RHSP to the SVNPs in their performance in regularizing cross loadings in Bayesian Regularized SEM.

In the remainder of this article we will first introduce regularization, whereby we will compare Bayesian and frequentist regularization. We will then outline Bayesian applications of regularized SEM. Here we will explain different shrinkage priors in detail, in particular the SVNPs and the RHSP. Next, a simulation study is reported that assesses the performance of the RHSP vs. the SVNPs in selecting cross loadings in a simple CFA model. We conclude by proposing directions for future research in further establishing the usefulness of the RHSP in Bayesian regularized SEM.

2 Theoretical Background

2.1 Regularization

A classic method of trying to find a balance between model complexity and model simplicity is *regularization* (Hastie, Tibshirani, & Wainwright, 2015). Regularization entails adding some bias to a model on purpose to reduce its variance. This helps to make models easier to interpret and more generalizable. In a frequentist context, regularization is achieved by adding a so-called penalty term to the cost function of a model. This ensures that model parameters that are irrelevant, e.g., small regression coefficients in a regression model with a large number of predictors, are shrunk to (or towards) zero. For a regression model, where we predict N individual scores y_i on an outcome y_i based on scores on a vector of predictors x_i , the vector of regression coefficients β and a random error term e_i :

$$y_i = \beta x_i + e_i, \text{ where}$$

$$e_i \sim \mathcal{N}(0, \sigma^2).$$

The Ordinary Least Squared Residuals estimates of β are obtained by minimizing the sum of squared residuals:

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \{ \sum_{i=1}^N (y_i - \beta x_i)^2 \}.$$

Penalized regression adds a a penalty term to this cost function, which is generally denoted as $\|\beta\|_L$:

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \{ \sum_{i=1}^N (y_i - \beta x_i)^2 + \lambda \|\beta\|_L \}.$$

When $L = 1$, we speak of the so-called L1 norm. In this case the penalty is:

$\|\beta\|_1 = \sum_{j=1}^p |\beta_j|$. This is the well-known LASSO (Least Absolute Shrinkage and Selection Operator) penalty (Tibshirani, 1996, 2011). Here, the absolute values of the regression coefficients are added up, multiplied by λ , and then added to the sum of the squared errors within the model's cost-function. The basic intuition is as follows. Just as minimizing the sum of the (squared) residuals leads to estimates of the model parameters with minimally small residuals, minimizing the (absolute) sum of the regression coefficients results in smaller values of the regression coefficients. The larger λ , the more weight the penalty has, and thereby the higher the amount of shrinkage to zero. In practice the hyperparameter λ is often determined by means of cross-validation. When $L = 2$ (the so-called L2 norm), $\|\beta\|_2 = \sqrt{\sum_{j=1}^p \beta_j^2}$. This is the famous ridge penalty (Hoerl & Kennard, 2000). Here, the same general principle is followed. This time not the absolute sum, but the square root of the sum of squares of the regression coefficients is minimized. In practice, one key difference between the LASSO and the ridge penalty is that the former shrinks some coefficient entirely to zero (thereby actively selecting predictors), whereas in ridge regression coefficients are only shrunken approximately but never entirely to zero.

Regularization can also be applied in more complex models, which is illustrated by the body of literature applying regularization in Structural Equation Modeling (SEM). Regularized SEM entails adding penalties to the cost function of SEM models (typically a variant of the maximum likelihood cost function, F_{ML}) to reach sparser models. Now, for

any given set of model parameters θ , we can add a penalty function $P(\theta)$ to the cost function of the model:

$$F_{ML, \text{penalized}} = F_{ML} + \lambda P(\theta).$$

Again the penalty function can take on different forms, such as the ridge (the square root of the sum of squares of the regularized parameters) or the LASSO (the sum of the absolute values of the regularized parameters), and λ is again a hyperparameter that determines the amount of shrinkage. Regularized SEM has successfully been applied in selecting cross loadings and residual covariances in CFA, especially under favorable conditions such as large sample sizes (Jacobucci, Grimm, & McArdle, 2016). Also structural model parameters, such as regression coefficients in MIMIC models (Jacobucci, Brandmaier, & Kievit, 2019; Jacobucci et al., 2016), or indirect effects in mediation models with continuous (Serang, Jacobucci, Brimhall, & Grimm, 2017) or dichotomous outcomes (Serang & Jacobucci, 2020), have successfully been regularized through the usage of penalties such as the LASSO or ridge penalty. R-packages such as *lsx* (Huang, 2020) or *regsem* (Li, Jacobucci, & Ammerman, 2021) have implemented classical penalties in various SEMs.

The key disadvantage of the frequentist regularization approach is that it depends on optimization. With more complicated penalties, and in particular for complex models, it can be hard to derive optimizable cost functions in practice. Also the derivation of unbiased standard errors is often challenging under such circumstances, which hinders reliable inference (Jacobucci & Grimm, 2018b; Jacobucci et al., 2016). Bayesian approaches to regularization overcome these obstacles (Jacobucci & Grimm, 2018b). In Bayesian model estimation, the so-called joint posterior distribution of the model parameters given the data is a combination of the data given the model parameters (the likelihood) and the prior distribution of the model parameters, i.e.,

$$P(\theta|data) \sim P(data|\theta)P(\theta).$$

Priors can therefore be used to shrink model parameter estimates to(wards) zero. Such

priors are called shrinkage priors. In the most simple case one can set a normal prior for parameters that is centered around zero (as the SVNP mentioned above), which resembles the ridge penalty (Hsiang, 1975). The LASSO penalty can be mimicked by setting a Laplace- (double exponential) prior for parameters (Hans, 2009; Park & Casella, 2008; see Van Erp et al., 2019 for the Bayesian equivalents of other relevant penalties). The fact that in Bayesian inference the whole joint posterior distribution of the model parameters can be sampled by means of Markov chain Monte Carlo (MCMC) methods poses a central advantage. This renders it unnecessary to derive standard errors, as the variance of the posterior distribution of model parameters is directly available. Also complex shrinkage priors with unique desired properties can be implemented without having to yield an optimizable cost-function, making Bayesian regularization a very flexible approach.

2.2 Bayesian CFA and The Small Variance Normal Prior (SVNP)

Confirmatory Factor Analysis (CFA, Bollen, 1989) is an essential tool for modeling measurement structures, falling under the class of Structural Equation Models (SEM). For every individual i , the scores on a vector of $p \times 1$ observed indicators y_i is modeled as follows:

$$y_i = \mu + \Lambda\eta_i + e_i,$$

where μ is a $p \times 1$ vector of intercepts, Λ is a $p \times q$ matrix of factor loadings, η_i is a $q \times 1$ vector of scores on the q latent factors, and e_i is a $p \times 1$ vector of random (measurement) error terms. Here, Λ is thus the part of the equation that relates the latent variables to the observed scores on the items.

We can differentiate between so-called main loadings, and cross loadings. The former are factor loadings that relate factors and items to one another that are theoretically expected to have a relationship. Cross loadings are factor loadings that relate factors to items that are not intended to measure that specific construct. One key property of CFA is that a clear expectation on the factor structure, i.e., which factor loads on which items, is

assumed. This allows to formulate clear expectations on which items can be viewed as cross loadings. Consequently, a natural way of identifying a confirmatory factor model (on top of other relevant identification constraints, such as scaling the factors) is to fix all cross loadings to zero. Not only does that identify the model, meaning that a unique solution can be found, but it also ensures that the measurement structure is easy to interpret. However, in practice, fixing all cross loadings to zero often results in poor model fit. So-called modification indices can be derived, which indicate how much model fit improves when adding or removing a certain parameter (for instance, a cross loading) to the model. Based on modification indices, researchers often re-add some cross loadings to the model, until a desired level of fit is reached. However, a core property of modification indices is that they are only based on the dataset at hand, and may therefore only pick up noise that is not relevant in the underlying population. Identifying cross loadings based on modification indices thus risks capitalization on chance. Consequently, measurement structures may be selected that do not generalize well to other datasets from the same population, hence over-fitting may occur (MacCallum et al., 1992).

As solution to this issue, Muthén and Asparouhov (2012), proposed a Bayesian way of identifying CFA models. They argued that fixing cross loadings of factor j on item k $\lambda_{c,jk}$ to zero can be interpreted as setting a normal prior with both a mean and a variance of zero on them:

$$\lambda_{c,jk} \sim \mathcal{N}(0, 0).$$

Now, since usually most cross loadings are indeed zero, but there are some that are a bit larger (in absolute terms) than zero, and few that are substantially larger, a more realistic prior would be a normal distribution that is still centered around zero, but does have a small variance, e.g.,

$$\lambda_{c,jk} \sim \mathcal{N}(0, 0.01).$$

The intuition of setting such a prior, which we refer to as Small Variance Normal Prior (SVNP), becomes clear when looking at Figure 1. In Bayesian model estimation the prior of a parameter directly influences the final posterior parameter estimate. By assuming that most cross loadings are zero, most posteriors of cross loadings will be centered around zero. At the same time, the prior allows for some deviations from zero, as there is also some prior mass for cross loadings that are marginally larger or smaller than zero. Of course, applying the SVNP as prior for cross loadings is thus nothing but a Bayesian variant of Regularized SEM, which strongly resembles the frequentist ridge penalty that was for instance applied by Jacobucci et al. (2016).

In practice an issue arises with the Bayesian CFA method, once substantially large cross loadings enter the picture. The estimates of truly large cross loadings should be large, as otherwise the (co-)variance structure of the data is not explained properly. With the SVNP, however, the estimates of large cross loadings are also shrunk substantially towards zero, as the prior with its thin tails attaches no prior mass to them. This causes bias. First, bias naturally occurs in the large cross loadings itself. However, also in other parameters, such as factor-correlations or main loadings, substantial bias can arise, as they are estimated conditionally on the cross loadings. To overcome this, the method is presented as a 2 step-approach (Muthén & Asparouhov, 2012). First, the model is estimated with the shrinkage prior set for the cross loadings. Cross loadings are then selected as non-zero if their 95% credible interval does not contain zero. Then the model is re-estimated. The as non-zero selected cross loadings are estimated freely without shrinkage, avoiding the bias in their estimates of the previous step. The as zero selected cross loadings are fixed to zero.

2.3 Alternative Shrinkage Priors

One alternative to the SVNP is the Bayesian LASSO, which has successfully been applied in regularized SEM previously (e.g. Chen, Guo, Zhihan, Zhang, Lijin, & Pan,

Junhao, 2021; Zhang et al., 2021). In general, its performance is similar to that of the SVNP. While a core difference between the ridge- and the LASSO penalty is that the LASSO penalty automatically shrinks some parameters to zero, the Bayesian equivalent of the LASSO also requires manual selection. This is because Bayesian model estimates based on posterior summaries, such as for instance the posterior means, will never be entirely zero (Zhang et al., 2021). As the SVNP, the Laplace prior of the Bayesian LASSO has thin tails, though slightly thicker than those of the SVNP, attaching only little prior mass to large cross loadings. Consequently, as with the SVNP, a two-step approach is required to estimate parameters without bias. Zhang et al. (2021) pointed out that it is difficult to formulate general recommendations on how to best select cross loadings as non-zero when applying the Bayesian LASSO. It is thus desirable to find shrinkage priors that can be applied in Bayesian CFA that allow for estimating model parameters without bias within a single step, to overcome the dependence on correctly selecting cross loadings.

One suitable alternative shrinkage prior for the purpose of selecting cross loadings in regularized Bayesian SEM is the so-called Spike-and-Slab Prior (George & McCulloch, 1993; Ishwaran & Rao, 2005; Mitchell & Beauchamp, 1988). This prior is a discrete mixture of an extremely peaked prior around zero (the spike), and a very flat prior for larger parameters (the slab). Formally, and applied to the cross loadings in CFA, for every cross loading of factor j on item k , the Spike-and-Slab Prior can be specified as (Lu et al., 2016):

$$\lambda_{c,jk}|r_{jk} \sim (1 - r_{jk})\delta_0 + r_{jk}\mathcal{N}(0, c_{jk}^2), \text{ with}$$

$$r_{jk} \sim \text{Bernoulli}(p_{jk}).$$

The basic intuition is as follows. When $r_{jk} = 1$, $\lambda_{c,jk} \sim \mathcal{N}(0, c_{jk}^2)$, hence $\lambda_{c,jk}$ is assigned to the slab. When $r_{jk} = 0$, $\lambda_{c,jk} \sim \delta_0$. In this case, the cross loading is thus assigned to a point-mass at zero, i.e., it is assigned to the spike. This ensures that large cross loadings that are relevant are not shrunk while small, negligible cross loadings are shrunk to zero. Lu et al. (2016) found that Spike-and-Slab Prior performs well in shrinking truly zero

cross loadings to zero, while not shrinking (relevant) large cross loadings to avoid bias, especially under favorable conditions with large sample sizes and cross loadings. The prior is thus in general more suited to estimate CFA models without substantial bias within a single estimation step, due to its refined nature in differentiating between large and small cross loadings. However, the Spike-and-Slab Prior cannot be implemented in Stan, one of the most popular software package for MCMC-sampling, as Stan does not allow for discrete mixture priors (Betancourt, 2018; Stan Development Team, 2021). This calls for a *non-discrete* alternative shrinkage prior that also outperforms the SVNP within a single estimation step.

2.4 The Regularized Horseshoe Prior (RHSP)

A fully continuous alternative to the Spike-and-Slab Prior that is implementable in Stan is the so-called *Regularized Horseshoe Prior* (RHSP, Piironen & Vehtari, 2017a, 2017b). This prior is an extension of the Horseshoe Prior (Carvalho, Polson, & Scott, 2010). The main idea of the original Horseshoe Prior is that there is a *global shrinkage parameter* τ , shrinking all cross loadings to zero. Next to this, there is a *local shrinkage parameter* $\bar{\omega}_{jk}$ ¹ that allows truly large cross loadings to escape the shrinkage, by setting thick Cauchy tails for the local scales ω_{jk} (Polson & Scott, 2010). Formally, the Horseshoe prior for every cross loading of factor j on item k is specified as follows:

$$\lambda_{c,jk} | \omega_{jk}, \tau, c \sim \mathcal{N}(0, \omega_{jk}^2 \tau^2), \text{ where}$$

$$\omega_{jk} \sim \mathcal{C}^+(0, 1).$$

The name-giving intuition behind the horseshoe prior becomes clear when considering the finding that a so-called shrinkage factor k_{jk} can be derived for the individual cross loadings (Carvalho et al., 2010; Piironen & Vehtari, 2017b). This shrinkage factor ranges from zero

¹ We deviate from the common notation of the local shrinkage parameter as $\bar{\lambda}$, as this letter is commonly used to denote factor loadings in CFA.

to one, with zero meaning no, and one meaning a lot of shrinkage. When plotting the density of k_{jk} there is a very high peak at low values and a very high peak at high values, resulting in a plot that resembles a horseshoe, illustrating that the Horseshoe Prior has the desired property of either shrinking parameters very little, or very much, with very few parameters that are shrunk in a non-extreme fashion.

The Horseshoe Prior was found consistently to possess the theoretical properties of not shrinking large parameters while shrinking small parameters substantially to zero in practice (Carvalho et al., 2010; Datta & Ghosh, 2013; Polson & Scott, 2010; Van Der Pas, Kleijn, & Van Der Vaart, 2014; Yamaguchi & Zhang, 2023). However, due to its heavy Cauchy tails it suffers from the same issues as a Cauchy prior. Specifically, not shrinking large parameters at all can lead to estimation issues, especially when parameters are weakly identified. This happens for instance in logistic regression with separable data, where a flat likelihood and thereby a weakly identified model arises (Ghosh, Li, & Mitra, 2018). The RHSP prevents such issues by shrinking also large parameters a little bit. The RHSP is specified as follows. For every cross loading of factor j on item k :

$$\begin{aligned}\lambda_{c,jk}|\bar{\omega}_{jk}, \tau, c &\sim \mathcal{N}(0, \bar{\omega}_{jk}^2 \tau^2), \text{ with } \bar{\omega}_{jk}^2 = \frac{c^2 \omega_{jk}^2}{c^2 + \tau^2 \omega_{jk}^2}, \\ \tau|df_{global}, s_{global} &\sim half - t_{df_{global}}(0, s_{global}^2), \text{ with } s_{global} = \frac{p_0}{p - p_0} \frac{\sigma}{\sqrt{N}}, \\ \omega_{jk}|df_{local}, s_{local} &\sim half - t_{df_{local}}(0, s_{local}^2), \\ c^2|df_{slab}, s_{slab} &\sim \mathcal{IG}(\frac{df_{slab}}{2}, df_{slab} \times \frac{s_{slab}^2}{2}),\end{aligned}$$

where p_0 represents a prior guess of the number of relevant cross loadings. Allowing to incorporate prior expectations on the expected number of non-zero parameters is another core advantage of the RHSP over the original Horseshoe Prior. However, it is not necessary to use p_0 . One can simply set s_{global} manually, whereby it is worth to consider that a s_{global} created based on a p_0 will typically be much lower than 1 (Piironen & Vehtari, 2017b).

Note that we specify the RHSP in its most general form. Setting the degrees of freedom of

the half-t-distributions to 1 results in half-Cauchy distributions. Strictly speaking, the prior is only a Regularized *Horseshoe* Prior when this is the case. In the current study we vary the degrees of freedoms of all scale parameters to assess the extent to which the sparcifying properties as well as the convergence of the RHSP are influenced by this.

The intuition of how the RHSP shrinks large parameters a little bit is best illustrated by assuming that c is a given constant. Now, when $\tau^2\omega_{jk}^2 < c^2$, $\bar{\omega}_{jk}^2 \rightarrow \omega_{jk}^2$. Hence, in this case the RHSP approaches the original Horseshoe Prior, with equally pronounced shrinkage to zero. However, when τ is far from zero, hence under large true cross loadings, $\tau^2\omega_{jk}^2 > c^2$, and $\bar{\omega}_{jk}^2 \rightarrow \frac{c^2}{\tau^2}$. Then, the prior of $\lambda_{c,jk}$ approaches a slab $\mathcal{N}(0, c^2)$. Under the above specification, when c is no constant but a parameter for which an Inverse-Gamma hyper-prior is set, the slab becomes a t-distribution with df_{slab} degrees of freedom, a mean of zero and a scale of s_{slab}^2 (Piironen & Vehtari, 2017b). Zhang and Liang (2023) showed that in MIMIC models, the RHSP indeed possesses the desired property of shrinking larger effects more than the regular Horseshoe Prior does, without negative consequences in terms of bias or parameter selection rates.

Figure 1 compares the two shrinkage priors that are the focus of our study. Both priors share a large peak at zero, which ensures that cross loadings are shrunken towards zero. However, the RHSP has much thicker tails. Here, for larger cross loadings, there is thus much more prior mass than with the SVNP. This ensures large cross loadings (and consequently other model parameters) can be estimated without bias within a single estimation step.

3 The current study

3.1 Study Procedure and Parameters

A Monte Carlo simulation study was conducted using Stan (Stan Development Team, 2021) and R (R Core Team, 2021). We used `cmdstandr` to interface Stan with R (Gabry &

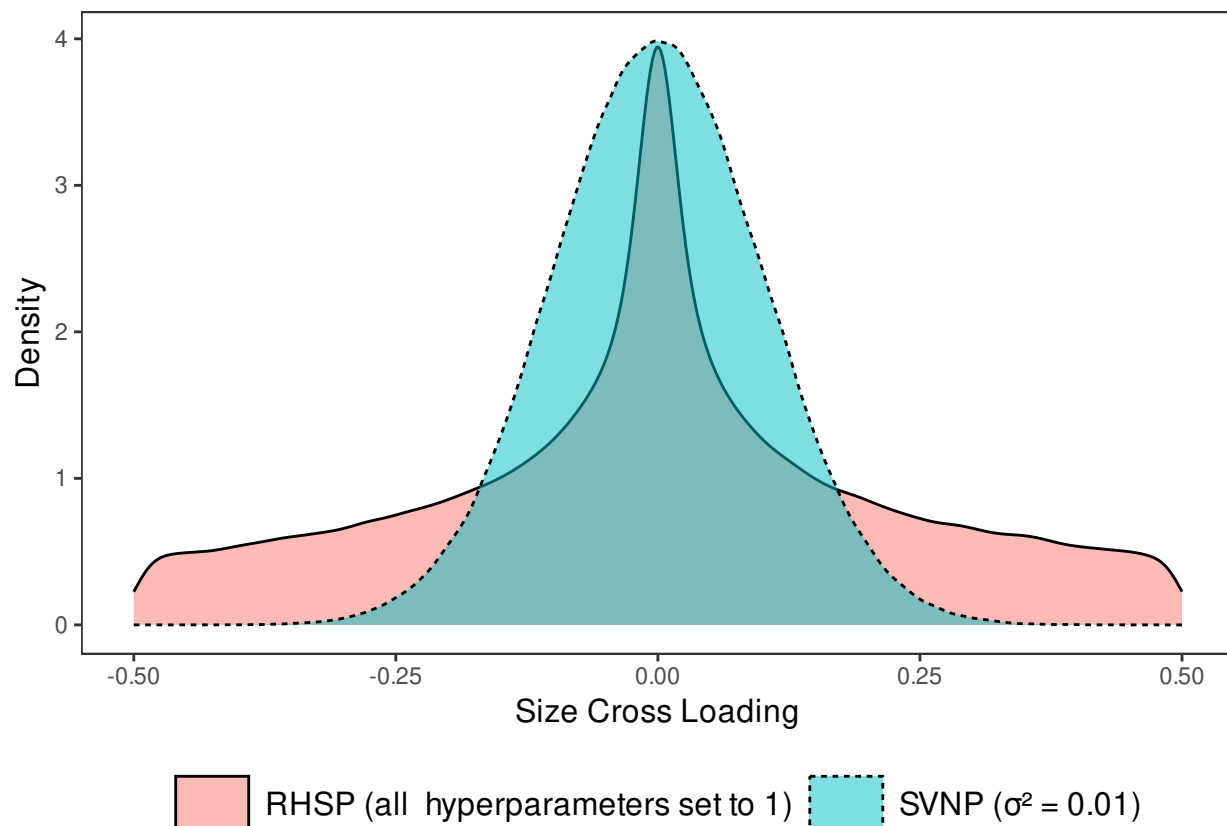


Figure 1. Density Plots of the Shrinkage Priors of Interest.

Češnovar, 2022), while also heavily relying on `rstan` (Stan Development Team, 2022) and `bayesplot` (Gabry & Mahr, 2022) for the post-processing of the posterior samples. Results were further post-processed using `tidyr` (Wickham & Girlich, 2022) & `dplyr` (Wickham, François, & Müller, 2022). All plots were made using `ggplot2` (Wickham, 2016). We ran up to 46 replications in parallel using the `parallel` package (R Core Team, 2022). This manuscript was written using R Markdown and the `papaja` package (Aust & Barth, 2022). All code that was used to run the simulations can be openly accessed on the author’s [github](#)². The models were sampled using the No-U-Turn-Sampler (Homan & Gelman,

² Specifically, the R-scripts needed to run the simulation can be found on <https://github.com/JMBKoch/1vs2StepBayesianRegSEM/tree/main/R>. `parameters.R` can be adjusted to adjust study parameters, and `main.R` is used to run the main simulation. Required packages are listed at the top of `parameters.R`.

2014), with two chains, a burnin-period of 2000 and a chain-length of 4000. These sampling parameters were identified in pilot runs to be required for the RHSP to reach convergence, and were therefore also used for the SVNP in order to ensure a fair comparison.

3.2 Conditions

3.2.1 Population Conditions. The datasets were simulated based on a true 2-factor model, with three items per factor, and a factor correlation of 0.5. The true model is summarized below, both in equations (Appendix A) and graphically (Figure 2).³ The factors were scaled by fixing their means to zero and their variances to 1. All main loadings were set to 0.75, and all residual variances to 0.3, to ensure that the largest proportion of variance in the items would be explained by their corresponding factor. We varied the size of the two truly non-zero cross loadings λ_{c1} and λ_{c6} between 0.2, a negligible magnitude such that shrinkage to zero is desired, and 0.5, a size for which shrinkage towards zero should be avoided. We varied the sample sizes of the simulated datasets between $N = 100$ and $N = 200$. Larger sample sizes of for instance 500 were not included despite being common place in previous simulation studies, because adding them would have rendered the run-time of the simulations for the RHSP unfeasible. This is appropriate because for simple factor models applied researchers are unlikely to collect such larger sample sizes in practice, such that our findings still generalize to real-life settings. In addition, as the sample size grows, the influence of the prior distribution decreases. Hence, the potential differences between the SVNP and RHSP will diminish with growing N .

For all parameters except the cross loadings, non-informative priors were set: for the main loadings a normal prior with mean of zero and a variance of 25, for the residual variances a half-Cauchy-prior with a location-parameter of 0 and a scale of 5, and for the factor-correlation Stan’s default uniform prior.

³ The Stan code of the model can be found on

<https://github.com/JMBKoch/1vs2StepBayesianRegSEM/blob/main/stan/>.

3.2.2 SVNP: Prior Conditions. We varied the hyperparameter of the SVNP, σ^2 , between 0.001, 0.01 and 0.1, based on Muthén and Asparouhov (2012). For the SVNP this left us with a total number of 2 (size cross loading) x 2 (N) x 3 (hyperparameter σ^2) = 12 individual conditions. Per condition, 200 replications were run, yielding a total of 2400 replications for this prior.

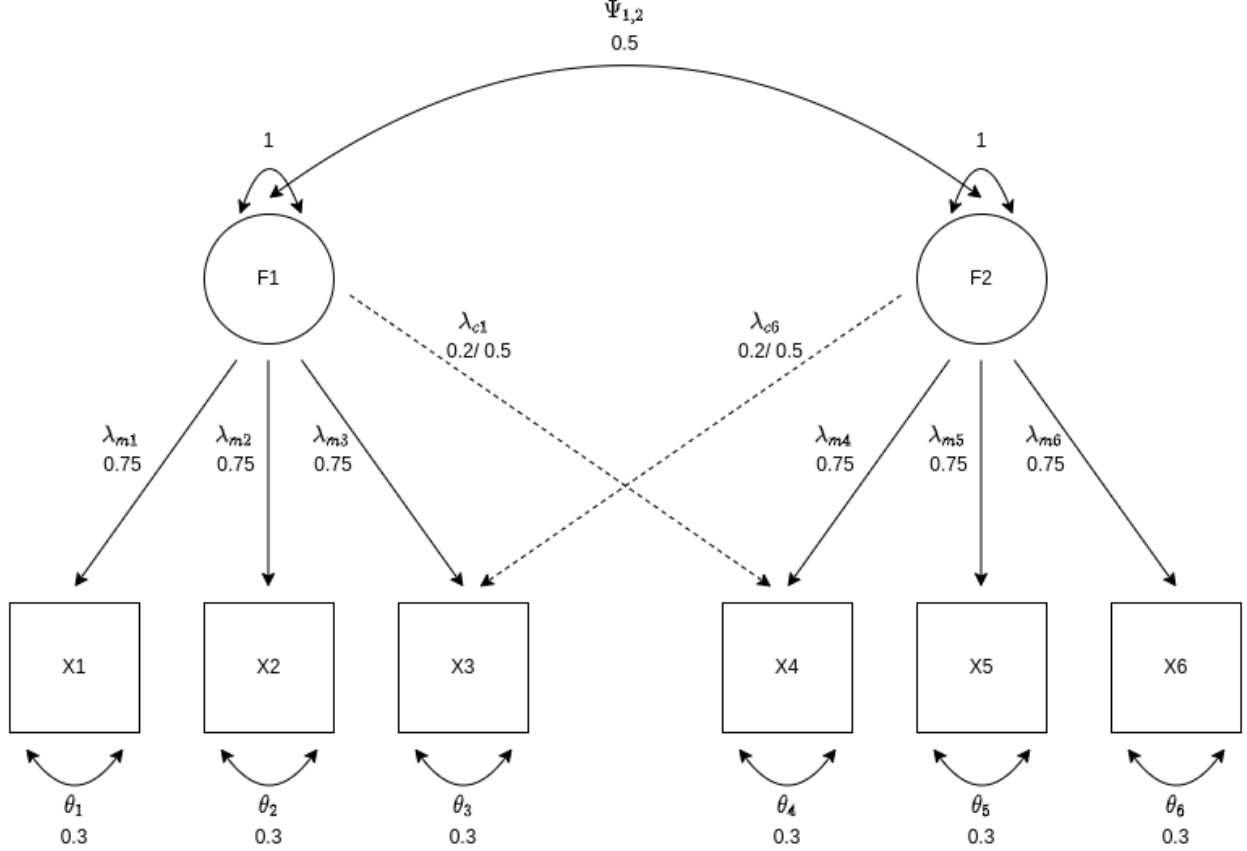


Figure 2. Graphical Representation of the True Model.

3.2.3 RHSP: Prior Conditions. The RHSP has six hyperparameters in the specification that we apply. We varied the scale of the global shrinkage parameter τ , s_{global} between, 0.1 and 1. Here 1, is a natural maximum given that the scale generally does not become larger than 1 when applying a prior guess p_0 (Piironen & Vehtari, 2017b), and 0.1 a logical minimum given the (standardized) scale of the model. Also the scale of the local shrinkage parameter ω_{jk} was varied between, 0.1 and 1. The degrees of freedoms of these two parameters, df_{local} and df_{global} were varied between 1 and 3. Larger degrees of

freedoms may help to overcome sampling issues that can arise when the df's are set to one. Finally, for the scale of the distribution of c^2 , s_{slab} was varied between 0.1, 1 and 5, and df_{slab} between 1 and 3. We decided to include a broader range of scales for the slab, as the slab is crucial in determining the shrinkage of large cross loadings. We were thus left with $2 (s_{global}) \times 2 (df_{global}) \times 2 (s_{local}) \times 2 (df_{local}) \times 3 (s_{slab}) \times 2 (df_{slab}) = 96$ hyperparameter conditions for the RHSP. In combination with the 2×2 population conditions this yielded 384 individual conditions for this prior. In total there were thus $384 \times 200 = 76800$ replications run for the RHSP. Table 1 summarizes all conditions.

3.3 Outcomes

All outcomes were computed based on both mean and median posterior estimates of the model parameters. We only present the results of the mean estimates, as the median estimates provided no relevant additional information⁴.

3.3.1 Mean Absolute Bias. For every model parameter θ and for every condition that has been sampled for N_{rep} replications, we computed the Mean Absolute Bias:

$$\bar{Bias}_{\bar{\theta}} = \frac{1}{N_{rep}} \sum_{i=1}^{N_{rep}} |\bar{\theta}_i - \theta_{true}|.$$

Given that the core issue of the SVNP is biased model estimates, this outcome naturally plays a central role in our study.

3.3.2 Relative Bias. The Relative Bias was computed per model parameter estimate and condition by dividing the estimates of the Mean Absolute Bias by the true value of the parameter:

$$Bias_{rel, \bar{\theta}} = \frac{\bar{Bias}_{\bar{\theta}}}{\theta_{true}}.$$

This outcome gives an indication of the magnitude of the bias by expressing it relative to the parameter's true value. However, given the standardized scale of the true model, the

⁴ see

<https://github.com/JMBKoch/1vs2StepBayesianRegSEM/tree/main/Rmd/analyses/medianEstimates>

Table 1

Overview study conditions.

Condition	Levels	Values
Population		
Sample size (N)	2	100, 200
Size cross loading $\lambda_{c1,6}$	2	0.2, 0.5
SVNP		
σ^2	3	0.1, 0.01, 0.001
RHSP		
s_{global}	2	0.1, 1
df_{global}	2	1, 3
s_{local}	2	0.1, 1
df_{local}	2	1, 3
s_{slab}	3	0.1, 1, 5
df_{slab}	2	1, 3

Note. SVNP = Small Variance Normal Prior.

RHSP = Regularized Horseshoe Prior. There was a total number of 396 conditions for both priors combined. All conditions were replicated 200 times.

Mean Absolute Bias is a quantity that can be interpreted rather intuitively in the context of this study and conclusions do not differ based on the relative bias.

3.3.3 Mean Squared Error: The Mean Squared Error (MSE) was computed per model parameter and condition as:

$$MSE_{\bar{\theta}} = \frac{1}{N_{rep}} \sum_{i=1}^{N_{rep}} (\bar{\theta}_i - \theta_{true})^2.$$

Another way to express the MSE is as the sum of the bias and the variance of a model parameter, which explains its added value over the Mean Absolute Bias alone. As with the Relative Bias the general pattern of results do not differ significantly for the MSE⁵.

3.3.4 Power and Type-I-Error Rate. Per condition, we computed the Power (true positive rate) in selecting truly non-zero cross loadings as non-zero by calculating the proportion of replications where the truly non-zero cross loadings were selected as non-zero. The Type-I-Error (false positive) rate in selecting truly zero cross loadings as non-zero was computed per condition as the proportion of replications where the truly zero cross loadings were selected as non-zero.

For both of these outcomes, we applied a variety of selection criteria for selecting cross loadings as non-zero, based on earlier research (Zhang et al., 2021). First, we used a variety of thresholding rules, where a cross loading is selected as non-zero when the absolute value of its estimate exceeds a specific threshold: 0, 0.05, 0.1, 0.15. Next, we considered four credible intervals (50%, 80%, 90%, 95%), where cross loadings are selected as non-zero when the interval does not contain zero.

4 Results

4.1 Convergence

4.1.1 Criteria. We removed replications for which at least one parameter did not converge. Convergence was determined based on two criteria. A parameter was viewed as

⁵ see <https://github.com/JMBKoch/1vs2StepBayesianRegSEM/tree/main/Rmd/analyses>

having reached convergence when its value of the effective sample size N_{Eff} exceeded 10% of the chain-length and its value of \hat{R} did not exceed 1.05. Moreover, we looked into the number and proportion of divergent transitions. Divergent transitions can be thought of as values of the Hamiltonian Monte Carlo Markov chain that diverge so much from their previous value that they cannot be trusted (see Betancourt, 2018 for details).

4.1.2 SVNP. The SVNP showed excellent performance in terms of convergence. Not a single replication had to be removed for not fulfilling the criteria outlined above. Moreover, across all runs there was not a single divergent transition. All 2400 replications of this prior were therefore included in the results.

4.1.3 RHSP. The RHSP showed weaker performance in terms of convergence than the SVNP. In total 742 replications were removed for not reaching convergence.

All replications of one condition were removed, because more than 50% (156 and thus 78%) of its replications failed entirely (i.e., not yielding any posterior samples) due to severe convergence problems. Specifically, this condition was:

$$N = 100, \text{ size } \lambda_{c1,6} = 0.2, s_{global} = s_{local} = s_{slab} = 0.1, df_{global} = df_{local} = df_{slab} = 1.$$

We removed the remaining 44 replications of this condition, as they were too little in number to give a reliable picture.

Next, we removed 542 replications across different conditions for not fulfilling the criteria outlined above. The maximum number of removed replications for a given condition was 37, which corresponds to 18.5% of the replications from this condition. Below in Table 2 we present all conditions under which more than 5% of the replications had to be removed. The conditions had in common that they had an N of 100, true cross loadings of 0.50, and a s_{global} and s_{local} of 0.1.

Table 3 presents the mean and maximum proportion of divergent transitions per condition, for all conditions under which there were, on average, at least 5% divergent transitions per chain. Again, these were the conditions that also suffered from convergence

Table 2

Conditions of the RHSP under which more than 5% of replications were removed for not having reached convergence ($N = 542$).

s_{global}	df_{global}	s_{local}	df_{local}	s_{slab}	df_{slab}	N	Size $\lambda_{c1,6}$	N removed Rep.
0.10	3	0.10	1	0.10	1	100	0.50	10
0.10	3	0.10	1	1.00	3	100	0.50	11
0.10	1	0.10	1	5.00	3	100	0.50	12
0.10	3	0.10	1	5.00	1	100	0.50	12
0.10	3	0.10	1	1.00	1	100	0.50	13
0.10	3	0.10	3	0.10	3	100	0.50	13
0.10	1	0.10	1	5.00	1	100	0.50	15
0.10	3	0.10	1	5.00	3	100	0.50	15
0.10	1	0.10	3	0.10	1	100	0.50	20
0.10	1	0.10	3	1.00	1	100	0.50	24
0.10	1	0.10	3	1.00	3	100	0.50	24
0.10	1	0.10	3	5.00	3	100	0.50	27
0.10	1	0.10	3	5.00	1	100	0.50	30
0.10	3	0.10	3	0.10	1	100	0.50	33
0.10	3	0.10	3	1.00	1	100	0.50	34
0.10	3	0.10	3	5.00	1	100	0.50	34
0.10	3	0.10	3	1.00	3	100	0.50	37
0.10	3	0.10	3	5.00	3	100	0.50	37

Note. Replications were removed for having an $\hat{R} \geq 1.05$ or an N_{eff} smaller than 10% of the chain-length, for any of the model parameters.

issues based on N_{Eff} and \hat{R} : $N = 100$, size cross loadings = 0.50, $s_{global} = s_{local} = 0.1$. We decided not to remove these replications, as this would have removed a substantial number of 2770 replications. In general, it is advised not to include any divergent transitions, since they introduce bias. Given the complex nature of the RHSP, which in practice usually leads to at least some divergent transitions, it is hard to follow this advise in practice. However, it needs to be taken into account in the interpretation of the findings that the divergent transitions may have added bias to the model estimates of the RHSP.

4.2 Main Results

4.2.1 Mean Absolute Bias. The Mean Absolute Bias of the SVNP and the RHSP for all parameters is summarized in Figure 3. For parameter estimates that show an identical pattern ($\bar{\lambda}_{c2-5}$, $\bar{\lambda}_{c1,6}$, $\bar{\lambda}_{m1,2,5,6}$, $\bar{\lambda}_{m3-4}$, and $\bar{\theta}_{1-6}$), the first respecting estimate is presented representative for all, both in Figure 3 and in the numbers presented below. As results are almost identical for the two sample sizes, we focus on presenting the findings for $N = 100$, to not distract from our main conclusions. Differences were so little that we do not present them here, to not distract from our main comparison to the SVNP. We decided to present the findings with all hyperparameters set to one, as this is a logical default hyperparameter configuration under the scale of a standardized CFA model. This is appropriate, as the results generally did not differ between different hyperparameter configurations. The replications for this hyperparameter configuration showed good convergence, such that only a single replication had to be removed.

The truly zero cross loadings ($\bar{\lambda}_{c2-5}$) had overall little levels of bias, across priors and hyperparameter settings. The patterns in the bias of the RHSP was almost identical to that of the SVNP with $\sigma^2 = 0.1$. Note that for both priors the bias in these cross loadings comes from these cross loading being under-estimated (i.e., estimated as smaller than zero).

Regarding the truly non-zero cross loadings ($\bar{\lambda}_{c1,6}$) the bias was sometimes very large

Table 3

Mean and Maximum proportion of divergent transitions averaged across replications and for conditions of the RHSP with on average more than 5% divergent transitions.

s_{global}	df_{global}	s_{local}	df_{local}	s_{slab}	df_{slab}	N	Size $\lambda_{c1,6}$	Mean	Max
0.10	1	0.10	3	0.10	1	100	0.50	0.06	0.56
0.10	1	0.10	3	5.00	1	100	0.50	0.06	0.71
0.10	1	0.10	3	5.00	3	100	0.50	0.06	0.75
0.10	3	0.10	1	0.10	1	100	0.50	0.07	0.66
0.10	3	0.10	1	5.00	1	100	0.50	0.06	0.58
0.10	3	0.10	3	0.10	1	100	0.50	0.07	0.54
0.10	3	0.10	3	1.00	1	100	0.50	0.07	0.65
0.10	3	0.10	3	5.00	1	100	0.50	0.07	0.83
0.10	3	0.10	3	5.00	3	100	0.50	0.08	0.69

Note. Mean = Mean proportion of divergent transitions, Max = maximum proportion of divergent transitions. There was a total of 2770 replications where the divergent transitions exceeded 5% of the chain-length. There were 14373 replications with more than 1% of divergent transitions. There were 1247 replications with more than 10% of divergent transitions. There were 46 replications with more than 50% of divergent transitions.

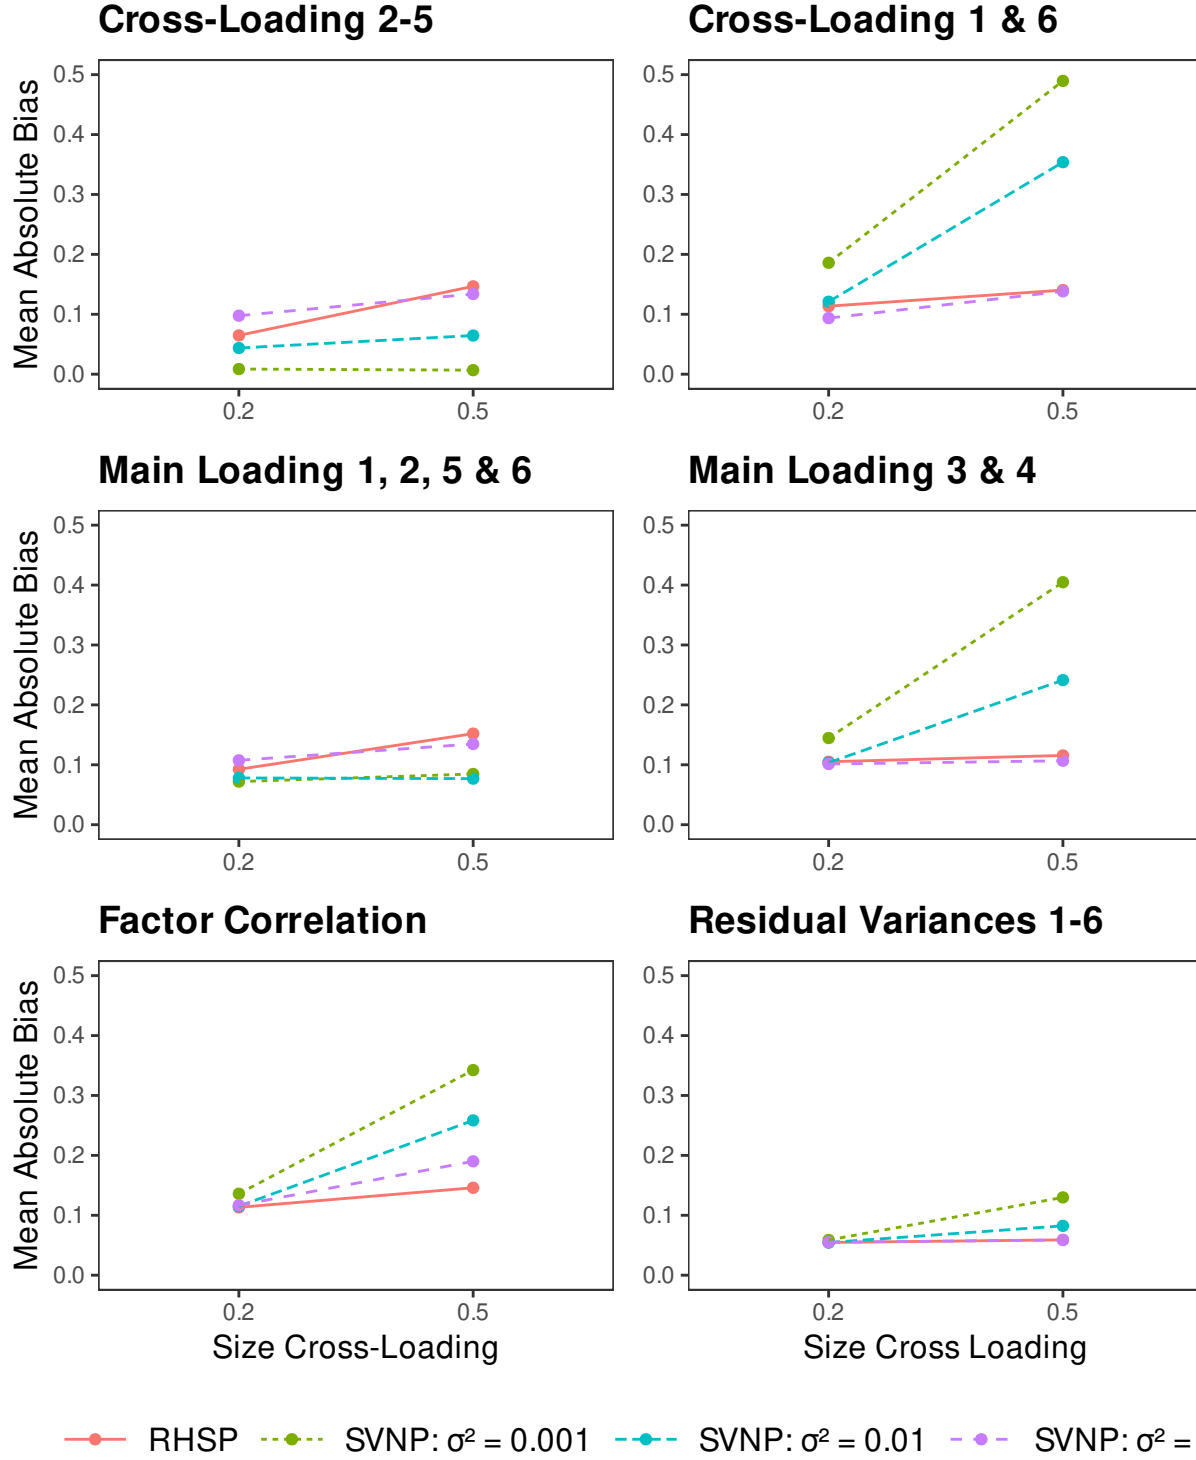


Figure 3. Mean Absolute Bias in the Model Parameters ($N = 100$). Per set of parameters that showed an identical pattern, the first parameter was used to represent all other parameters, e.g., cross loading 2 was plotted representative for cross loading 3-5. All hyperparameters of the RHSP are set to 1 in the results presented here.

for the SVNPN, particularly with large true cross loadings of 0.5 and $\sigma^2 = 0.001$ (e.g., $Bias_{\bar{\lambda}_{c1}} = 0.49$), since the estimates of the true cross loadings of 0.5 were shrunk almost entirely to zero (e.g., $\bar{\lambda}_{c1} = 0.01$). Also with $\sigma^2 = 0.01$ (and true cross loadings of 0.5) substantial bias occurred (e.g., $Bias_{\bar{\lambda}_{c1}} = 0.35$), as the cross loadings were still under-estimated considerably (e.g., $\bar{\lambda}_{c1} = 0.15$), though not entirely shrunk to zero. With $\sigma^2 = 0.1$ the bias in the estimates of the truly non-zero cross loadings of 0.5 was less pronounced (e.g., $Bias_{\bar{\lambda}_{c1}} = 0.14$). Here σ^2 was large enough to estimate the cross loadings closer to their true value, e.g., $\bar{\lambda}_{c1} = 0.37$. The bias of the RHSP was again almost entirely identical to that of the SVNPN with $\sigma^2 = 0.1$.

Also the estimates of the main loadings of factor 1 on item 3 ($\bar{\lambda}_{m3}$) and of factor 2 on item 4 ($\bar{\lambda}_{m4}$) were substantially biased under the SVNPN when the true cross loadings were 0.5 and $\sigma^2 = 0.001$ (e.g., $Bias_{\bar{\lambda}_{m3}} = 0.40$). These two loadings showed much higher bias than the other four main loadings, as they loaded on the same two items as the two non-zero cross loadings ($\bar{\lambda}_{c1}$ and $\bar{\lambda}_{c6}$, see Figure 2). As the cross loadings were shrunk to zero, these main loadings now also accounted for the variance in the items that was truly explained by the cross loadings. Consequently, the two main loadings were over-estimated, e.g., $\bar{\lambda}_{m3} = 1.15$. Both the RHSP and the SVNPN with $\sigma^2 = 0.1$ had low levels of bias for the main loadings both with true cross loadings of 0.2 and 0.5.

In the factor correlation $\bar{\Psi}_{1,2}$ the bias was relatively small and approximately the same for the RHSP and the SVNPN with different values of σ^2 when the truly non-zero cross loadings were 0.2. For the SVNPN, bias became again much more pronounced with true cross loadings of 0.5 and small values of σ^2 , particularly $\sigma^2 = 0.001$ ($Bias_{\bar{\Psi}_{1,2}} = 0.34$). In this situation the factor correlation was heavily over-estimated ($\bar{\Psi}_{1,2} = 0.84$). This can be explained by the fact that the covariance between item 3 and 4 that arose from the two cross loadings, was misattributed to the factor-correlation, as the cross loadings were shrunk to zero. For both the SVNPN with $\sigma^2 = 0.1$ and the RHSP, the bias in the factor

correlation was much lower, although there was still a noticeable increase between true cross loadings of 0.2 and 0.5.

For both priors, the bias in the estimates of the residual variances $\bar{\theta}_{1-6}$ was not large across different conditions, although for the SVNP there was a noticeable increase between true cross loadings of 0.2 and 0.5 when $\sigma^2 = 0.001$.

4.2.2 Power and Type-I-Error Rate. Figure 4 summarizes the Power and the Type-I-Error rate of both priors. Again, the outcomes are presented for the first parameter of an identical set of parameters (e.g., $\lambda_{c,1}$ is presented representative for the two truly non-zero cross loadings). Moreover, results of the RHSP are again presented with all hyperparameters set to one. In the plots summarizing the Power the horizontal red dashed line indicates the minimum power of .80 recommended by Muthén and Asparouhov (2012). According to Cham, West, Ma, and Aiken (2012), the maximally acceptable Type-I-Error rate is the upper-bound of a 95% interval of a binomial distribution, in this case $0.05 + 1.96 \times \sqrt{0.05 \times (1 - 0.05)/N_{rep}} = 0.08$. This maximum is visualized by the horizontal red dashed line in the Type-I-Error rate plots. For both priors, with a threshold of 0.00 there is both a Power and Type-I-Error rate of 1. This is logical, since posterior means will never be entirely zero (Zhang et al., 2021). The result thus mostly serves to illustrate this property of Bayesian inference and thereby the need for alternative selection rules in Bayesian regularization, if the goal is variable selection itself, and not only unbiased model parameter estimates. Whether a high Power or a low Type-I-Error rate is prioritized depends on many factors, such as the model at hand and the research goal. In general, one should aim for a balance between both, with high enough levels of Power and low enough Type-I-Error rates co-occurring.

For both priors and across most conditions, the Power fell under the desired threshold of .80. For the SVNP with $\sigma^2 = 0.001$ non-zero cross loadings were always over-shrunk so much that they were never selected as non-zero. Under $\sigma^2 = 0.01$, the situation




Figure 4. Mean Power and Type-I-Error Rates in Selecting non-zero Crossloadings. All hyperparameters of the RHSP are set to 1 in the results presented here.

improved somewhat, with now cross loadings of 0.5 being correctly selected as non-zero for all selection rules when $N = 200$. For $N = 100$, thresholds of 0.05, and 50% credible intervals also had the desired levels of power, with thresholds of 0.10 also almost reaching a power of .80. The SVNPs performed best in terms of power when $\sigma^2 = 0.1$. With $N = 200$, all selection rules except for the 95% credible intervals reached the desired Power. With $N = 100$, all selection rules except for the 95% and 90% credible intervals exceeded a power of .80. For the RHSP, in general the Power did also not live up to the standards by Muthén and Asparouhov (2012). The desired power was exclusively reached in selecting cross loadings that were truly 0.5, when using 50% credible intervals or the three thresholds. Especially the widest (90% and 95%) credible intervals performed very poorly, which is explained by the fact that the lower bounds of these intervals (almost) always exceeded zero.

In general, under most conditions the Type-I-Error rate of both the SVNPs and the RHSP exceeded the desired maximum. For the SVNPs, with $\sigma^2 = 0.001$ and $N = 100$, the Type-I-Error rate stayed very low for all selection criteria (up to the 0.00 threshold), as under this condition all cross loadings, including the truly zero ones, were always shrunk almost entirely to zero. With $N = 200$, the 50% credible intervals lead to an undesired large Type-I error rate. With $\sigma^2 = 0.01$ and $N = 100$, some selection rules (90% & 95% credible intervals, a threshold of 0.15) had an acceptable Type-I-Error rate even with large true cross loadings of 0.5, but they quickly became unacceptably high with $N = 200$. With $\sigma^2 = 0.1$, all selection rules except for 90% and 95% credible intervals had unacceptably high Type-I-Error rates for both sizes of non-zero cross loadings, even though pronounced differences between cross loadings of 0.2 and 0.5 occurred. This was the condition under which the SVNPs had the best levels of power, indicating that the good performance in Power comes with the caveat of a poor Type-I-Error rate. The RHSP performed even worse than the SVNPs in terms of its Type-I-Error rate. Only the 90% and 95% credible intervals stayed within the desired boundary, as they always included zero. This, however,

is not useful in practice as these intervals were completely unable to identify truly non-zero cross loadings correctly as non-zero.

5 Conclusions and Discussion

This was the first study to apply the Regularized Horseshoe Prior (RHSP, Piironen & Vehtari, 2017b) to select cross loadings in CFA. A comparison to the Bayesian CFA approach by Muthén and Asparouhov (2012) was made, where cross loadings are regularized with the Small Variance Normal Prior (SVNP).

Both the SVNP and the RHSP performed well with small truly non-zero cross loadings of 0.2, in terms of estimating the model without substantial bias. This can be interpreted as a successful instance of regularization, where an acceptable amount of bias is added to the model by shrinking some parameters to(wards) zero, to reach a sparser solution. However, under some hyperparameter settings the performance of the SVNP decreased substantially with larger truly non-zero cross loadings of 0.5. With smaller values of σ^2 , particularly with $\sigma^2 = 0.001$, these cross loadings were still shrunk to zero, even though they were much larger in reality. This caused substantial bias, not only in the estimates of the cross loadings itself, but also in the estimates of some main loadings and the factor correlation. In practice, bias in structural parameters is particularly concerning, as it may lead to wrong conclusions in research on structural relationships between latent constructs. With $\sigma^2 = 0.1$, the SVNP was able to estimate the model parameters with little bias. Such relatively large variance still allowed for enough deviations from zero in the cross loadings to yield relatively accurate estimates of the non-zero cross loadings itself and, consequently, the other model parameters. The RHSP performed very similar to the SVNP with $\sigma^2 = 0.1$.

While the RHSP and the SVNP were both able to estimate the simple CFA model of this study with little bias, only the RHSP was able to do so consistently across different

hyperparameter settings. This poses a central advantage over the SVNPs. Being robust to different hyperparameter settings, which is achieved through the complex nature of the RHSP that sets hyper-priors on all parameters, may ultimately allow researchers to apply the RHSP without having to worry about choosing optimal values for the hyperparameters. However, the other side of the coin is that the complex nature of the RHSP makes it sensitive to hyperparameter settings regarding its performance in terms of convergence, in contrast to the SVNPs that had excellent convergence across all conditions. Hereby, a clear pattern arose, with convergence issues arising with large cross loadings, small sample sizes and small scales for the local and global shrinkage parameter. While this may suggest that one should simply avoid using small values for the scales, using larger scales that shrink parameters less may lead to identification problems in more complex SEM models. Future research is required that incorporates more complex models, in order to assess the performance of the RHSP in terms of convergence in more depth. Ultimately, it is desirable to be able to formulate clear guidelines on which hyperparameter configurations to use in which situations to avoid convergence issues, to make the RHSP a practical approach.

Particularly under simple models such as the one employed in this study using the less complex SVNPs with a relatively large value of σ^2 may thus prove advantageous in practice, to avoid risking bias stemming from non-convergence and divergent transitions. Next to this, the simpler SVNPs are easier to implement and take substantially less time to run. However, simply using larger values of σ^2 with the SVNPs is no general solution. In practice, models may include more structural parameters, even more cross loadings, or a number of residual co-variances. Under these circumstances, large values of σ^2 may lead to identification issues. Moreover, the larger σ^2 , the more cross loadings will be selected as non-zero, which may ultimately lead to over-fitting. Because of this, the RHSP might prove as more advisable for more complex models. While the RHSP, as the SVNPs, can be assumed to run into more and more identification issues with increasing model complexity, its refined nature in giving prior mass only where strictly desired (see Figure 1) might allow

for a more efficient identification of complex models than the SVNPs, at least for some hyperparameter settings. Using the RHSP may thus turn out to be the approach of choice in situations where identification is hard to achieve with the SVNPs. However, future research comparing the SVNPs to the RHSP in regularizing more complex SEM-models is yet to investigate this directly.

Regarding the Power and Type-I-Error rate in selecting cross loadings as non-zero, both priors performed poorly across a range of selection rules. First of all, this is not surprising, given earlier research which clearly showed that a range of shrinkage priors (SVNP, Bayesian LASSO, Spike-and-Slab Prior, Lu et al., 2016; Zhang et al., 2021) generally need much larger sample sizes than the ones employed in this study to reach desirable levels of Power and Type-I-Error rates (see Jacobucci et al., 2016; Lu et al., 2016, who show that frequentist variable selection methods are sensitive to sample size too). At the same time, it is an important finding that both priors perform poorly under small sample sizes, as simple CFA models are often fitted with such small sample sizes in practice. Interestingly, Zhang and Liang (2023) found the RHSP to perform well in selecting covariates in a MIMIC model as (non-)zero, even under a sample size of $N = 200$. This suggests that the small sample sizes in our study may not be the only reason for the poor selection performance of the RHSP. This is yet to be explored systematically in future research, by assessing the Power and Type-I-Error rate of the RHSP under larger sample sizes (e.g., $N = 500, 1000, 2000$). Note, however, that our findings do not imply that the RHSP is useless under small sample sizes. If the goal of regularization is not to select which cross loadings are non-zero, but to yield unbiased estimates of the other model parameters, the RHSP still works better than the SVNPs under most settings, even with small sample sizes. In practice, researchers often fit SEM-models including a measurement structure to test structural hypotheses. For this purpose, the question of whether or not cross loadings are zero is not as relevant, as long as the structural model parameters are estimated without substantial bias. Next to this, different selection methods, such as

Highest Posterior Density (HPD) intervals (Zhang et al., 2021), or projection predictive inference (Piironen & Vehtari, 2017b) may be more successful in making a selection that meets the desired balance in Power and Type-I-Error rate.

The large levels of bias found in the SVNP (for the majority of hyperparameter settings) are not surprising, given the original approach explicitly asking for a second step to circumvent bias (Muthén & Asparouhov, 2012). However, the fact that the SVNP performed so poorly in selecting the cross loadings across selection rules suggests that in practice the 2-step approach, which heavily relies on a correct selection of cross loadings as non-zero, is not advisable. This, however, needs to be directly assessed in future research. One approach of doing so would be to first make a selection of cross loadings and then actively execute the second estimation step using new validation data sets. This would allow to compare the bias of the 2-step SVNP to the (1-step) RHSP directly. While our study formed a valuable first step in establishing the performance of the RHSP in regularized SEM, such a design would allow to more directly answer the question of whether or not the RHSP enables us to ‘get a step ahead’.

Next to the specific limitations named above, the current study has some general shortcomings. First, we only assessed the performance of the RHSP in regularizing cross loadings in a very simple CFA model consisting of only two factors. A straightforward way of extending the current study and making its findings more generalizable would be to include factor models with more factors. Within CFA models, another important set of parameters that can be identified through the usage of shrinkage priors are residual co-variances (i.e., the off-diagonal elements of Θ), which are usually fixed to zero in classic CFA (Muthén & Asparouhov, 2012). An important next step is thus to assess the performance of the RHSP in selecting residual co-variances on top of cross loadings. Hereby, it is advisable to start by regularizing residual co-variances separately rather than on top of cross loadings, as regularizing both parameters simultaneously can be challenging

in practice (Zhang et al., 2021). It is also desirable to assess the performance of the RHSP in regularizing model parameters in more complex SEM models. For instance, extending the work by Zhang and Liang (2023) who regularized the causes (covariates) of latent variables, but not their indicators (i.e. factor loadings), may be a sensible step. Hereby, adding another latent variable and with it cross loading to their model would shed more light onto the properties of the RHSP in regularizing both latent and structural parameters simultaneously. Also the regularization of other structural parameters, such indirect effects in mediation models is an interesting research direction. Moreover, applying the RHSP in measurement models with non-continuous (i.e., binary, ordinal or nominal) outcomes would be an interesting way of building on the current study. Regarding the SVNP, incorporating extensions of the SVNP where a hyper-prior is set for σ^2 (see for instance Lu et al., 2016) may prove valuable in allowing for a more refined comparison between the SVNP and the RHSP. Note also that in the current study no direct comparison to other relevant shrinkage priors, such as the Bayesian LASSO or the Spike-And-Slab Prior was made. Van Erp (2023) emphasizes the need for user-friendly software for Bayesian Regularized SEM. In particular for more complex shrinkage priors software implementations are lacking. This also counts for the RSHP, which needs to be implemented into standard Bayesian SEM software (e.g., by adding it to the BLAVAAN package, Merkle, Fitzsimmons, Uanhoro, & Goodrich, 2020). Only then applied researchers can actually use it in practice. Here, a particular challenge that the RHSP has such a large number of hyperparameters which may differ in results based on the data and model at hand. For this however, the desired clear guidelines on hyperparameter settings for the RHSP mentioned above need to be derived first. Only then user-friendly software with sensible default-settings can be designed.

Despite the limitations named, the current study forms a valuable contribution to the current literature. We set a crucial first step in establishing the usefulness of the RHSP in Bayesian regularized CFA. Our findings show that the RHSP can generally be successfully be applied in Bayesian regularized SEM and point to important directions for future

research.

References

- Aust, F., & Barth, M. (2022). *Papaja: Prepare reproducible APA journal articles with R Markdown*. Retrieved from <https://github.com/crsh/papaja>
- Betancourt, M. (2018). A Conceptual Introduction to Hamiltonian Monte Carlo. *arXiv:1701.02434 [Stat]*. Retrieved from <http://arxiv.org/abs/1701.02434>
- Bollen, K. A. (1989). *Structural Equations with Latent Variables*. John Wiley & Sons.
- Carvalho, C. M., Polson, N. G., & Scott, J. G. (2010). The horseshoe estimator for sparse signals. *Biometrika*, 97(2), 465–480.
<https://doi.org/10.1093/biomet/asq017>
- Cham, H., West, S. G., Ma, Y., & Aiken, L. S. (2012). Estimating Latent Variable Interactions With Nonnormal Observed Data: A Comparison of Four Approaches. *Multivariate Behavioral Research*, 47(6), 840–876.
<https://doi.org/10.1080/00273171.2012.732901>
- Chen, J., Guo, Zhihan, Zhang, Lijin, & Pan, Junhao. (2021). A Partially Confirmatory Approach to Scale Development With the Bayesian Lasso. *Psychological Methods*, 26(2), 210–235. Retrieved from <https://oce-ovid-com.proxy.library.uu.nl/article/00060744-202104000-00005/HTML>
- Cox, D. R. (2006). *Principles of Statistical Inference*. Cambridge University Press.
- Datta, J., & Ghosh, J. K. (2013). Asymptotic properties of Bayes risk for the horseshoe prior. *Bayesian Analysis*, 8(1), 111–132.
- Gabry, J., & Češnovar, R. (2022). *Cmdstanr*. Retrieved from <https://mc-stan.org/cmdstanr/>
- Gabry, J., & Mahr, T. (2022). *Bayesplot: Plotting for Bayesian Models*. Retrieved from <https://mc-stan.org/bayesplot/>
- George, E. I., & McCulloch, R. E. (1993). Variable Selection Via Gibbs Sampling. *Journal of the American Statistical Association*, 88(423), 881–889.

- <https://doi.org/10.2307/2290777>
- Ghosh, J., Li, Y., & Mitra, R. (2018). On the Use of Cauchy Prior Distributions for Bayesian Logistic Regression. *Bayesian Analysis*, 13(2), 359–383.
- <https://doi.org/10.1214/17-BA1051>
- Hans, C. (2009). Bayesian lasso regression. *Biometrika*, 96(4), 835–845. Retrieved from <https://www.jstor.org/stable/27798870>
- Hastie, T., Tibshirani, R., & Wainwright, M. (2015). Statistical learning with sparsity. *Monographs on Statistics and Applied Probability*, 143, 143.
- Hoerl, A. E., & Kennard, R. W. (2000). Ridge Regression: Biased Estimation for Nonorthogonal Problems. *Technometrics*, 42(1), 80–86.
- <https://doi.org/10.2307/1271436>
- Homan, M. D., & Gelman, A. (2014). The No-U-turn sampler: Adaptively setting path lengths in Hamiltonian Monte Carlo. *The Journal of Machine Learning Research*, 15(1), 1593–1623.
- Hsiang, T. C. (1975). A Bayesian View on Ridge Regression. *Journal of the Royal Statistical Society. Series D (The Statistician)*, 24(4), 267–268.
- <https://doi.org/10.2307/2987923>
- Huang, P.-H. (2020). lsx: Semi-confirmatory structural equation modeling via penalized likelihood. *Journal of Statistical Software*, 93(7), 1–37.
- <https://doi.org/10.18637/jss.v093.i07>
- Ishwaran, H., & Rao, J. S. (2005). Spike and slab variable selection: Frequentist and Bayesian strategies. *The Annals of Statistics*, 33(2), 730–773.
- <https://doi.org/10.1214/009053604000001147>
- Jacobucci, R., Brandmaier, A. M., & Kievit, R. A. (2019). A Practical Guide to Variable Selection in Structural Equation Modeling by Using Regularized Multiple-Indicators, Multiple-Causes Models. *Advances in Methods and Practices in Psychological Science*, 2(1), 55–76.

- <https://doi.org/10.1177/2515245919826527>
- Jacobucci, R., & Grimm, K. J. (2018b). Comparison of Frequentist and Bayesian Regularization in Structural Equation Modeling. *Structural Equation Modeling: A Multidisciplinary Journal*, 25(4), 639–649.
- <https://doi.org/10.1080/10705511.2017.1410822>
- Jacobucci, R., & Grimm, K. J. (2018a). Comparison of frequentist and bayesian regularization in structural equation modeling. *Structural Equation Modeling: A Multidisciplinary Journal*, 25(4), 639–649.
- Jacobucci, R., Grimm, K. J., & McArdle, J. J. (2016). Regularized Structural Equation Modeling. *Structural Equation Modeling: A Multidisciplinary Journal*, 23(4), 555–566. <https://doi.org/10.1080/10705511.2016.1154793>
- James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). *An Introduction to Statistical Learning: With Applications in R*. New York, NY: Springer US.
- <https://doi.org/10.1007/978-1-0716-1418-1>
- Li, X., Jacobucci, R., & Ammerman, B. A. (2021). Tutorial on the use of the regsem package in r. *Psych*, 3(4), 579–592.
- Lu, Z.-H., Chow, S.-M., & Loken, E. (2016). Bayesian Factor Analysis as a Variable-Selection Problem: Alternative Priors and Consequences. *Multivariate Behavioral Research*, 51(4), 519–539.
- <https://doi.org/10.1080/00273171.2016.1168279>
- MacCallum, R. C., Roznowski, M., & Necowitz, L. B. (1992). Model modifications in covariance structure analysis: The problem of capitalization on chance. *Psychological Bulletin*, 111(3), 490–504.
- <https://doi.org/10.1037/0033-2909.111.3.490>
- Merkle, E. C., Fitzsimmons, E., Uanhoro, J., & Goodrich, B. (2020). Efficient Bayesian Structural Equation Modeling in Stan. *arXiv:2008.07733 [Stat]*. Retrieved from <http://arxiv.org/abs/2008.07733>

- Mitchell, T. J., & Beauchamp, J. J. (1988). Bayesian Variable Selection in Linear Regression. *Journal of the American Statistical Association*, 83(404), 1023–1032.
<https://doi.org/10.2307/2290129>
- Muthén, B., & Asparouhov, T. (2012). *Bayesian SEM: A more flexible representation of substantive theory*. 78. <https://doi.org/10.1037/a0026802>
- Park, T., & Casella, G. (2008). The Bayesian Lasso. *Journal of the American Statistical Association*, 103(482), 681–686.
<https://doi.org/10.1198/0162145080000000337>
- Piironen, J., & Vehtari, A. (2017a). On the Hyperprior Choice for the Global Shrinkage Parameter in the Horseshoe Prior. *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, 905–913. PMLR. Retrieved from <https://proceedings.mlr.press/v54/piironen17a.html>
- Piironen, J., & Vehtari, A. (2017b). Sparsity information and regularization in the horseshoe and other shrinkage priors. *Electronic Journal of Statistics*, 11(2), 5018–5051. <https://doi.org/10.1214/17-EJS1337SI>
- Polson, N. G., & Scott, J. G. (2010). Shrink globally, act locally: Sparse Bayesian regularization and prediction. *Bayesian Statistics*, 9(501-538), 105.
- R Core Team. (2021). *R: A Language and Environment for Statistical Computing*. Retrieved from <https://www.R-project.org/>
- R Core Team. (2022). *Package 'parallel'*. Retrieved from <https://stat.ethz.ch/R-manual/R-devel/library/parallel/doc/parallel.pdf>
- Serang, S., & Jacobucci, R. (2020). Exploratory Mediation Analysis of Dichotomous Outcomes via Regularization. *Multivariate Behavioral Research*, 55(1), 69–86.
<https://doi.org/10.1080/00273171.2019.1608145>
- Serang, S., Jacobucci, R., Brimhall, K. C., & Grimm, K. J. (2017). Exploratory Mediation Analysis via Regularization. *Structural Equation Modeling: A Multidisciplinary Journal*, 24(5), 733–744.

- <https://doi.org/10.1080/10705511.2017.1311775>
- Stan Development Team. (2021). *Stan User Guide*. Retrieved from https://mc-stan.org/docs/2_27/stan-users-guide-2_27.pdf
- Stan Development Team. (2022). *Rstan: The R interface to Stan*. Retrieved from <https://mc-stan.org/>
- Tibshirani, R. (1996). Regression Shrinkage and Selection Via the Lasso. *Journal of the Royal Statistical Society: Series B (Methodological)*, 58(1), 267–288. <https://doi.org/10.1111/j.2517-6161.1996.tb02080.x>
- Tibshirani, R. (2011). Regression shrinkage and selection via the lasso: A retrospective. *Journal of the Royal Statistical Society. Series B (Statistical Methodology)*, 73(3), 273–282. Retrieved from <https://www.jstor.org/stable/41262671>
- Van Der Pas, S. L., Kleijn, B. J., & Van Der Vaart, A. W. (2014). The horseshoe estimator: Posterior concentration around nearly black vectors. *Electronic Journal of Statistics*, 8(2), 2585–2618.
- Van Erp, S. (2023). Bayesian regularized SEM: Current capabilities and constraints. *Psych*, 5(3), 814–835.
- Van Erp, S., Oberski, D. L., & Mulder, J. (2019). Shrinkage priors for Bayesian penalized regression. *Journal of Mathematical Psychology*, 89, 31–50. <https://doi.org/10.1016/j.jmp.2018.12.004>
- Wickham, H. (2016). *ggplot2: Elegant Graphics for Data Analysis*. Retrieved from <https://ggplot2.tidyverse.org>
- Wickham, H., François, R., & Müller, K. (2022). *Dplyr: A Grammar of Data Manipulation*. Retrieved from <https://dplyr.tidyverse.org>, <https://github.com/tidyverse/dplyr>
- Wickham, H., & Girlich, M. (2022). *Tidyr: Tidy Messy Data*. Retrieved from <https://tidyr.tidyverse.org>, <https://github.com/tidyverse/tidyr>

- Yamaguchi, K., & Zhang, J. (2023). Fully gibbs sampling algorithms for bayesian variable selection in latent regression models. *Journal of Educational Measurement*, 60(2), 202–234.
- Zhang, L., & Liang, X. (2023). Bayesian regularization in multiple-indicators multiple-causes models. *Psychological Methods*.
- Zhang, L., Pan, J., & Ip, E. H. (2021). Criteria for Parameter Identification in Bayesian Lasso Methods for Covariance Analysis: Comparing Rules for Thresholding, p -value, and Credible Interval. *Structural Equation Modeling: A Multidisciplinary Journal*, 1–10. <https://doi.org/10.1080/10705511.2021.1945456>

5.1 Appendix A: True Model

For every individual i in $i = 1, \dots, N$:

$$Y_i \sim \mathcal{N}(\mathbf{0}, \Sigma),$$

where

$$\begin{aligned} \Sigma &= \Lambda \Psi \Lambda', \\ \Lambda &= \begin{bmatrix} 0.75 & 0 \\ 0.75 & 0 \\ 0.75 & 0.2/0.5 \\ 0.2/0.5 & 0.75 \\ 0 & 0.75 \\ 0 & 0.75 \end{bmatrix}, \\ \Psi &= \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix}, \end{aligned}$$

and

$$\Theta = \text{diag}[0.3, 0.3, 0.3, 0.3, 0.3, 0.3].$$